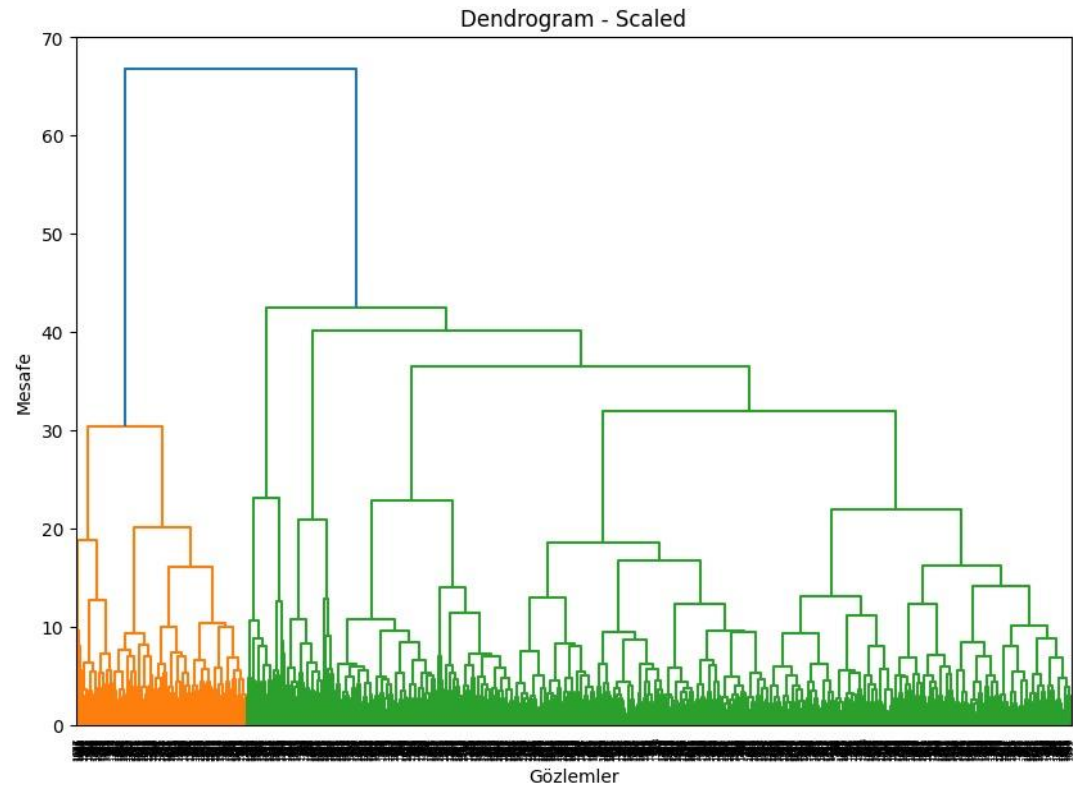
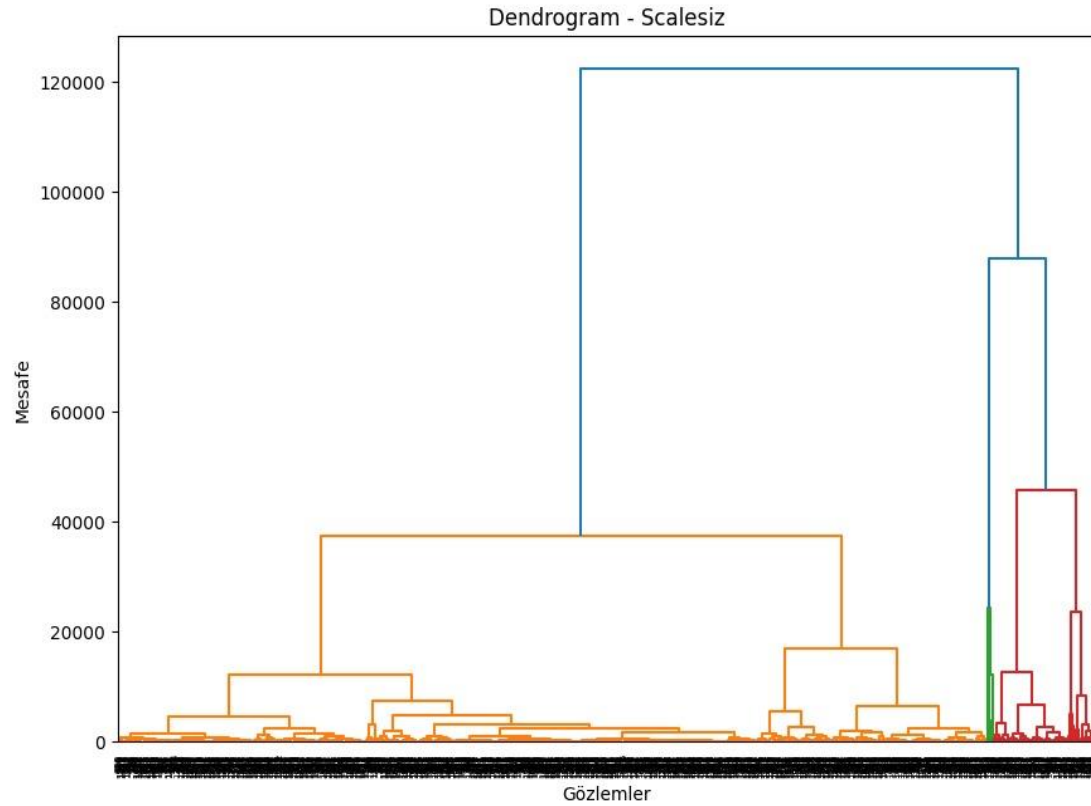


# Efficient Telemarketing

Enes YILDIRIM 20211603001

Demir SARRAÇ 20231603050

# Dendrograms




# Is Data Set Balanced?

Unfortunately, our dataset contains 4000 "no" and 521 "yes" labels.

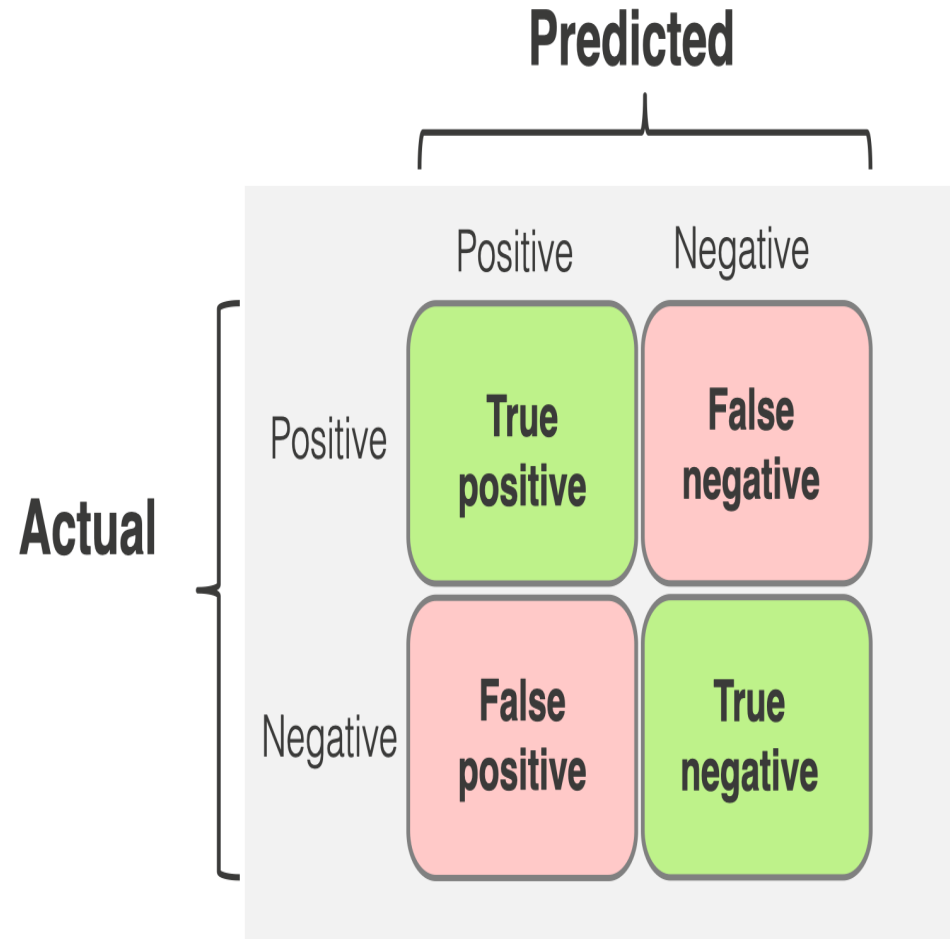
This means our dataset is imbalanced.

We will balance the dataset by downsampling the majority class.

Later, to compare, we will apply SMOTE.



- Recall that, using our dataset, we are trying to predict whether the client will subscribe to a deposit or not.
- Since we have **limited human resources**, we aim to predict those who will subscribe in advance to increase efficiency.
- To do this, we **prioritize the recall metric**, as it helps us minimize false negatives and ensures we don't miss potential subscribers.
- In this case, a **false negative** means **missing a potential customer** who would actually subscribe — leading to a significant loss in efficiency and potential revenue.



```

no_df = df[df["y"] == "no"]
yes_df = df[df["y"] == "yes"]

downsampled_no = resample(no_df,
                           replace=False,
                           n_samples = len(yes_df),
                           random_state=57
)

balanced_df = pd.concat([yes_df, downsampled_no])
print(balanced_df.info())
print(balanced_df.shape)
print(balanced_df["y"].value_counts())

```

✓ 0.0s

```

... <class 'pandas.core.frame.DataFrame'>
Index: 1042 entries, 13 to 979
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1042 non-null   int64
1   job         1042 non-null   object
2   marital     1042 non-null   object
3   education   1042 non-null   object
4   default     1042 non-null   object
5   balance     1042 non-null   int64
6   housing     1042 non-null   object
7   loan        1042 non-null   object
8   contact     1042 non-null   object
9   day         1042 non-null   int64
10  month       1042 non-null   object
11  duration    1042 non-null   int64
12  campaign    1042 non-null   int64
13  pdays      1042 non-null   int64
14  previous    1042 non-null   int64
15  poutcome    1042 non-null   object
16  y           1042 non-null   object
dtypes: int64(7), object(10)
memory usage: 146.5+ KB
None
...
y
yes    521
no     521
Name: count, dtype: int64

```

# Logistic Regression & Treshold

```
logistic_model = LogisticRegression(random_state=57)
logistic_model.fit(X_train_scaled, y_train)

y_pred = logistic_model.predict(X_test_scaled)

print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

✓ 0.0s

```
[[69 35]
 [37 68]]
```

	precision	recall	f1-score	support
0	0.65	0.66	0.66	104
1	0.66	0.65	0.65	105
accuracy			0.66	209
macro avg	0.66	0.66	0.66	209
weighted avg	0.66	0.66	0.66	209

```
y_probs = logistic_model.predict_proba(X_test_scaled)[: ,1]
treshold = 0.4

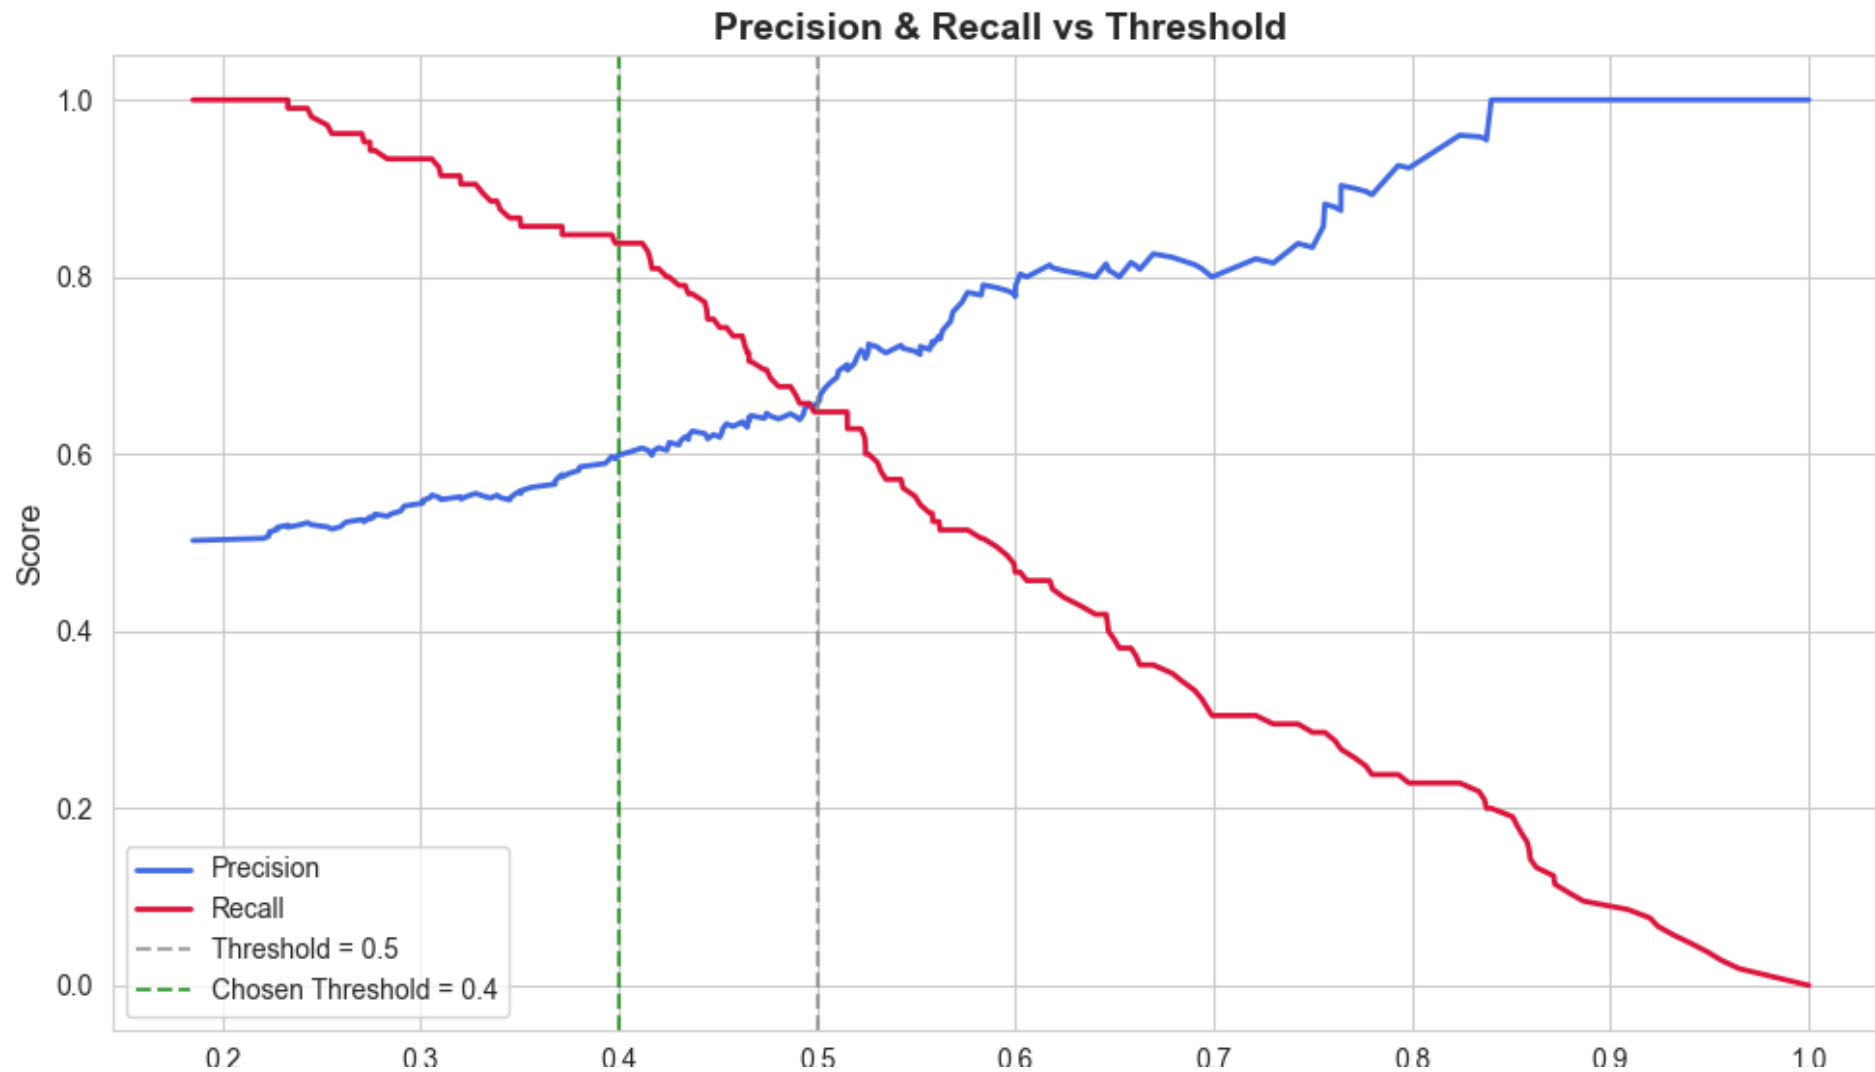
y_pred_custom = (y_probs >= treshold).astype(int)

print(confusion_matrix(y_test, y_pred_custom))
print(classification_report(y_test, y_pred_custom))
```

✓ 0.0s

```
[[46 58]
 [17 88]]
```

	precision	recall	f1-score	support
0	0.73	0.44	0.55	104
1	0.60	0.84	0.70	105
accuracy			0.64	209
macro avg	0.67	0.64	0.63	209
weighted avg	0.67	0.64	0.63	209



- To find the best trade-off point, we need to know the cost of each type of mistake.
- However, a detailed cost-sensitive analysis is beyond the scope of our current project

# KNN & Neighbour Number

```
knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train_scaled, y_train)

y_pred = knn.predict(X_test_scaled)

print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

✓ 0.0s

```
[[67 37]
 [42 63]]
```

	precision	recall	f1-score	support
0	0.61	0.64	0.63	104
1	0.63	0.60	0.61	105
accuracy			0.62	209
macro avg	0.62	0.62	0.62	209
weighted avg	0.62	0.62	0.62	209

```
knn = KNeighborsClassifier(n_neighbors=7)
knn.fit(X_train_scaled, y_train)

y_pred = knn.predict(X_test_scaled)

print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

✓ 0.0s

```
[[69 35]
 [47 58]]
```

	precision	recall	f1-score	support
0	0.59	0.66	0.63	104
1	0.62	0.55	0.59	105
accuracy			0.61	209
macro avg	0.61	0.61	0.61	209
weighted avg	0.61	0.61	0.61	209



# Decision Tree Classifier & Random Forest Classifier

```
dt_model = DecisionTreeClassifier(random_state=57, max_depth=5)

dt_model.fit(X_train_scaled, y_train)

y_pred = dt_model.predict(X_test_scaled)

print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

✓ 0.0s

```
[[33 71]
 [17 88]]
```

	precision	recall	f1-score	support
0	0.66	0.32	0.43	104
1	0.55	0.84	0.67	105
accuracy			0.58	209
macro avg	0.61	0.58	0.55	209
weighted avg	0.61	0.58	0.55	209

```
rf_model = RandomForestClassifier(n_estimators=100, max_depth=5, random_state=57)
rf_model.fit(X_train_scaled, y_train)
y_pred = rf_model.predict(X_test_scaled)

print(rf_model.score(X_test_scaled, y_test))
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))



cm = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = cm.ravel()
total = tn + fp + fn + tp
```

✓ 0.0s

0.6746411483253588

```
[[83 21]
 [47 58]]
```

	precision	recall	f1-score	support
0	0.64	0.80	0.71	104
1	0.73	0.55	0.63	105
accuracy			0.67	209
macro avg	0.69	0.68	0.67	209
weighted avg	0.69	0.67	0.67	209



logistic regression scores 0.6514193781815121 0.631492673992674

random forest scores 0.7304332886565799 0.6555494505494506

# Logistic Regression With SMOTE

```
logistic_model = LogisticRegression(random_state=57)
logistic_model.fit(X_train_res_scaled, y_train_res)

y_pred = logistic_model.predict(X_test_scaled)

print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

✓ 0.0s

```
[[745  56]
 [ 82  22]]
```

	precision	recall	f1-score	support
0	0.90	0.93	0.92	801
1	0.28	0.21	0.24	104
accuracy			0.85	905
macro avg	0.59	0.57	0.58	905
weighted avg	0.83	0.85	0.84	905

```
y_probs = logistic_model.predict_proba(X_test_scaled)[: ,1]
threshold = 0.4

y_pred_custom = (y_probs >= threshold).astype(int)

print(confusion_matrix(y_test, y_pred_custom))
print(classification_report(y_test, y_pred_custom))
```

✓ 0.0s

```
[[716  85]
 [ 70  34]]
```

	precision	recall	f1-score	support
0	0.91	0.89	0.90	801
1	0.29	0.33	0.30	104
accuracy			0.83	905
macro avg	0.60	0.61	0.60	905
weighted avg	0.84	0.83	0.83	905

In the graphs above, SMOTE was applied and we used a balanced dataset consisting of 4000 "yes" and 4000 "no" labels. Interestingly, downsampling performed better than SMOTE in terms of recall. Additionally, we once again observed that lowering the threshold increases recall.